

Process Synchronization

- Background
- The Critical-Section Problem
- Semaphores
- Classical Problems of Synchronization
- Monitors

Background

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating/concurrent processes.
- That mechanism/logic is called process synchronization.



The Critical-Section Problem

- n processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.
- Structure of process P_i

repeat


entry section

critical section

exit section

reminder section

until false;



Solution to Critical-Section Problem

1. **Mutual Exclusion.** If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.

If one process is accessing the shared data other should wait.

2. **Progress.** All the concurrent process involved in the synchronization (mutual exclusion) must progress and driven to completion.
3. **Bounded Waiting.** There should be definite waiting time for processes following mutual exclusion.

Semaphore

- It is the solution to critical section problem.
- Helps in achieving mutual exclusion.
- Semaphore S – integer variable
- can only be accessed via two indivisible (atomic) operations

wait (S): **while** $S \leq 0$ **do** *no-op*;
 $S := S - 1$;

signal (S): $S := S + 1$;

Example: Critical Section of n Processes

- Shared variables
 - **var** *mutex* : *semaphore*
 - initially *mutex* = 1
- Process P_i

repeat

wait(mutex);

0: locked

critical section

signal(mutex);

1: opened

remainder section

until *false*;

Semaphore Implementation

- Define a semaphore as a record

type *semaphore* = **record**

value: integer

L: list of process;

end;

- Assume two simple operations:
 - Block: suspends the process that invokes it.
 - wakeup(*P*): resumes the execution of a blocked process *P*.

Implementation (Cont.)

- Semaphore operations now defined as

wait(S): $S.value := S.value - 1;$

if $S.value < 0$

then begin

 add this process to $S.L$;
 block;

end;

signal(S): $S.value := S.value + 1;$

if $S.value \leq 0$

then begin

 remove a process P from $S.L$;
 wakeup(P);

end;

Semaphore as General Synchronization Tool

- Execute B in P_j only after A executed in P_i
- Use semaphore $flag$ initialized to 0
- Code:

P_i
 \vdots
 A
 $signal(flag)$

P_j
 \vdots
 $wait(flag)$
 B

B can not be executed
until unless process P_i
execute the $signal(flag)$
instruction

$wait(S)$: **while** $S \leq 0$ **do** *no-op*;
 $S := S - 1$;
 $signal(S)$: $S := S + 1$;

Deadlock and Starvation

- Deadlock – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let S and Q be two semaphores initialized to 1

P_0	P_1
$wait(S);$	$wait(Q);$
$wait(Q);$	$wait(S);$
\vdots	\vdots
$signal(S);$	$signal(Q);$
$signal(Q)$	$signal(S);$

- Starvation – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Types of Semaphores

- *Counting semaphore* – integer value can range over an unrestricted domain.
- *Binary semaphore* – integer value can range only between 0 and 1; can be simpler to implement.